

THREADS

22UCSC401 - Object Oriented Programming

4th Semester, B Division, Academic Year: 2023 - 24

Department of Computer Science and Engineering, SDMCET, Dharwad-2

THREADS

- ◉ Java provides support for Multithreaded programming
- ◉ A multithreaded program contains two or more parts that can run **concurrently**
- ◉ Each part of such program is called as **thread**
- ◉ Each thread defines a **separate path of execution**
- ◉ Multithreading is a special form of multitasking

THREADS

- ◉ Two types of multi-tasking:
 - Process-based
 - Thread-based
- ◉ Process-based multi-tasking allows the computer to run two or more programs concurrently
- ◉ Thread-based multi-tasking allows computer to run two or more threads concurrently
- ◉ In case of thread-based multi-tasking, **thread is the smallest unit of dispatch-able code**

THREADS

Threads	Processes
Light-weight	Heavy-weight
Less overhead	High overhead
Inter-thread communication is inexpensive	Inter-process communication is expensive
Context switching is inexpensive	Context switching is expensive
Share the same address space	Have separate address space

JAVA THREAD MODEL

- ◉ Single-threaded systems use an approach called an **event loop** with **polling**
- ◉ In such systems, when a thread *blocks* (that is, suspends execution) because it is waiting for some resource, the entire program stops running
- ◉ In multi-threading systems, event loop and polling is **eliminated**

JAVA THREAD MODEL

- ⦿ One thread can pause without stopping other parts of the program
- ⦿ When a thread blocks in a Java program, only the single thread that is blocked pauses
- ⦿ All other threads continue to run

THREAD STATES

- ◉ **Running** (a thread can be running)
- ◉ **Ready** (a thread can be ready to run as soon as it gets CPU time)
- ◉ **Suspend** (a thread can be suspended temporarily from its activity)
- ◉ **Resume** (only a suspended thread can be resumed)
- ◉ **Block** (a thread can be blocked when waiting for a resource)

THREAD PRIORITIES

- ◉ Java assigns to each thread a priority that determines how that thread should be treated with respect to the others
- ◉ Thread priorities are integers that specify the relative priority of one thread over another
- ◉ A thread's priority is used to decide when to switch from one running thread to the next
- ◉ This is called **context switch**

THREAD PRIORITIES

- ◉ Rules that determine when a context switch takes place:
 - A thread can **voluntarily relinquish** control
 - This is done by explicitly **yielding, sleeping, or blocking on pending I/O**
 - In this scenario, all other threads are examined, and the highest-priority thread that is ready to run is given the CPU
 - A thread can be **preempted** by a higher-priority thread
 - A lower-priority thread that does not yield the processor is simply **preempted—no matter what it is doing**—by a higher-priority thread
 - Basically, as soon as a higher-priority thread wants to run, it does
 - This is called **preemptive multitasking**

THREAD PRIORITIES

- ◉ In cases where two threads with the same priority are competing for CPU cycles, the situation is a bit complicated
- ◉ For Windows, threads of equal priority are time-sliced automatically in round-robin fashion
- ◉ For other OSes', threads of equal priority must voluntarily yield control to their peers
- ◉ If they don't, the other threads will not run

SYNCHRONIZATION

- ◉ Multithreading introduces an **asynchronous behavior** to the programs
- ◉ So there must be a way to enforce synchronicity when we want
- ◉ This is taken care in Java by the use of “**monitors**”
- ◉ Monitor mechanism was defined by C.A.R.Hoare

SYNCHRONIZATION

- ◉ We can think of monitors as a very small box that can hold only one thread
- ◉ Once a thread enters a monitor, all other threads must wait until that thread exits the monitor
- ◉ Monitors can be used to protect a shared data from being manipulated by multiple threads at a time

SYNCHRONIZATION

- ◉ In Java, there is no class called “monitor”
- ◉ Instead, each object has its own implicit monitor that is automatically entered when one of the object’s synchronized methods is called
- ◉ Once a thread is inside a synchronized method, no other thread can call any other synchronized method on the same object

THREAD CLASS & RUNNABLE INTERFACE

- ◉ Java's multithreading system is built upon the *Thread* class, its methods, and its companion interface, *Runnable*
- ◉ To create a new thread, Java program will either extend *Thread* or implement the *Runnable* interface

METHODS OF THREAD CLASS

Method	Meaning
getName	Obtain a thread's name.
getPriority	Obtain a thread's priority.
isAlive	Determine if a thread is still running.
join	Wait for a thread to terminate.
run	Entry point for the thread.
sleep	Suspend a thread for a period of time.
start	Start a thread by calling its run method.

THE MAIN THREAD

- ◉ When a Java program starts up, one thread begins running immediately
- ◉ This is usually called the **main thread**
- ◉ The main thread is important for two reasons:
 - It is the thread from which other “child” threads will be spawned
 - Often, it must be the last thread to finish execution because it performs various shutdown actions

THE MAIN THREAD

- ⦿ Although the main thread is created automatically when the program is started, it can be controlled through a *Thread* object
- ⦿ To do this, we must first obtain a reference to it by calling the method *currentThread()*, which is a *public static* member of *Thread*

THE MAIN THREAD

- ⦿ General form of *currentThread()*:
 - *static Thread currentThread()*
- ⦿ This method returns a reference to the thread in which it is called
- ⦿ Once we have a reference to the main thread, we can control it just like any other thread

THE MAIN THREAD

- ◉ The method *sleep()* causes the thread from which it is called to suspend its execution for the specified period of milliseconds
- ◉ General form:
 - *static void sleep(long milliseconds) throws InterruptedException*
 - *static void sleep(long milliseconds, int nanoseconds) throws InterruptedException*

THE MAIN THREAD

- ◉ General form of *setName()* method:

- *final void setName(String threadName)*

- ◉ General form of *getName()* method:

- *final String getName()*

CREATING A THREAD

- ⦿ There are two ways in Java to create a thread:
 - Implement *Runnable* interface
 - Extend *Thread* class

CREATING A THREAD BY IMPLEMENTING RUNNABLE INTERFACE

- ◉ Create a class that implements *Runnable* interface
- ◉ To implement *Runnable* interface, a class need only to implement a single method called *run()*
 - *public void run()*
- ◉ In *run()*, we need to write the code that constitutes the new thread

CREATING A THREAD BY IMPLEMENTING RUNNABLE INTERFACE

- ◉ *run()* can call other methods, use other classes, and declare variables
- ◉ *run()* establishes an entry point for the thread
- ◉ When *run()* returns, the thread ends its execution

CREATING A THREAD BY IMPLEMENTING RUNNABLE INTERFACE

- ◉ After a class implements *Runnable* interface, it also has to instantiate an object of type *Thread* from within that class
- ◉ **Thread** has several constructors
- ◉ The one that is used here is:
 - *Thread(Runnable threadObj, String threadName)*
- ◉ *threadObj* is an instance of a class that implements *Runnable* interface
- ◉ Name of the new thread is specified by *threadName*

CREATING A THREAD BY IMPLEMENTING RUNNABLE INTERFACE

- ◉ After the new thread is created, we have to start its execution by calling *start()* method, which is declared within *Thread*
- ◉ *start()* executes a call to *run()*
- ◉ General form of *start()*:
 - *void start()*

CREATING A THREAD BY INHERITING THREAD CLASS

- ◉ Create a new class and make it to extend *Thread*
- ◉ Then create an instance of that class
- ◉ The extending class must override the *run()* method, which acts as the entry point for the new thread
- ◉ The extended class must also call *start()* to begin execution of new thread

CREATING A THREAD BY INHERITING THREAD CLASS

- ◉ Child thread was created by instantiating an object of *NewThread* which is extended from *Thread*
- ◉ Call to *super()* in the program invokes the following form of *Thread* constructor:
 - *public Thread(String threadName)*
- ◉ *threadName* specifies name of the thread

CHOOSING AN APPROACH

- ◉ Which approach to use for creating threads in Java??

USING ISALIVE() AND JOIN()

- ⦿ How can one thread know when another thread has ended?
 - *final boolean isAlive()*
- ⦿ The *isAlive()* method returns true if the thread upon which it is called is still running
- ⦿ It returns *false* otherwise

USING ISALIVE() AND JOIN()

- ◉ If we want to wait for a thread to finish its execution, use *join()* method:
 - *final void join() throws InterruptedException*
- ◉ Waits until the thread on which it is called terminates

THREAD PRIORITIES

- ◉ To set a thread's priority, use the *setPriority()* method, which is a member of *Thread*
 - *final void setPriority(int level)*
- ◉ The value of *level* must be within the range *MIN_PRIORITY* and *MAX_PRIORITY*
- ◉ *MIN_PRIORITY=1* & *MAX_PRIORITY=10*
- ◉ *NORM_PRIORITY=5* (default priority)

THREAD PRIORITIES

- ◉ To obtain the current priority, use the *getPriority()* method of *Thread*
 - *final int getPriority()*

SYNCHRONIZATION

- ◉ A Monitor is an object that is used as mutually exclusive lock, or **mutex**
- ◉ Only one thread can own monitor at a time
- ◉ When a thread acquires a lock, it is said to have entered the monitor
- ◉ All the other threads attempting to enter the locked monitor will be suspended until the first thread exits the monitor

SYNCHRONIZATION

- ⦿ Concept of synchronization is implemented in Java by two ways:
 - Using synchronized methods
 - Using synchronized statement
- ⦿ Both involve the use of *synchronized* keyword

USING SYNCHRONIZED METHODS

- ◉ In Java, all objects have their own implicit monitor associated with them
- ◉ To enter an object's monitor, just call a method that has been modified with the *synchronized* keyword
- ◉ When a thread is inside a synchronized method, all other threads that try to call on the same instance have to wait
- ◉ To exit the monitor, the owner of the monitor simply returns the synchronized method

```
class Display {  
    void output(String msg) {  
        System.out.print("[ " + msg);  
        try {  
            Thread.sleep(1000);  
        } catch (InterruptedException e) {  
            e.printStackTrace();  
        }  
        System.out.println("]");  
    }  
}
```

```
class MyThreads implements Runnable {  
    private String msg;  
    Thread t;  
    Display d;  
    MyThreads(Display d, String msg) {  
        this.d = d;  
        this.msg = msg;  
        t = new Thread(this);  
        t.start();  
    }  
    public void run() {  
        d.output(msg);  
    }  
}
```

```
public class SynchronizationDemo {  
    public static void main(String[] args) {  
        Display d = new Display();  
        MyThreads mt1 = new MyThreads(d, "SDM");  
        MyThreads mt2 = new MyThreads(d, "CET");  
        MyThreads mt3 = new MyThreads(d, "CSE");  
        try {  
            mt1.t.join();  
            mt2.t.join();  
            mt3.t.join();  
        } catch (InterruptedException e) {  
            e.printStackTrace();  
        }  
    }  
}
```

USING SYNCHRONIZED STATEMENT

- General form of synchronized statement:

```
synchronized(object) {  
    // statements to be synchronized  
}
```

INTER-THREAD COMMUNICATION

- ◉ To implement ITC, Java uses *wait()*, *notify()* & *notifyAll()* methods and these are implemented as *final* methods in *Object*
- ◉ All these methods can only be called from within a *synchronized* context

INTER-THREAD COMMUNICATION

- ◉ *wait()* tells the calling thread to give up the monitor and go to sleep until some other thread enters the same monitor and calls *notify()*.
- ◉ *notify()* wakes up a thread that called *wait()* on the same object.
- ◉ *notifyAll()* wakes up all the threads that called *wait()* on the same object. One of the threads will be granted access.

INTER-THREAD COMMUNICATION

- ⦿ These three methods are declared within *Object*:
- ⦿ *final void wait() throws InterruptedException*
- ⦿ *final void notify()*
- ⦿ *final void notifyAll()*

PRODUCER-CONSUMER PROBLEM

- ⦿ Contains 4 classes:

- Q, the queue that needs to be synchronized
- Producer, the threaded object that is producing queue entries
- Consumer, the threaded object that is consuming queue entries
- PC, the tiny class that creates a single Q, Producer and Consumer

```
class Q {  
    int n;  
    synchronized int get() {  
        System.out.println("Got: " + n);  
        return n;  
    }  
    synchronized void put(int n) {  
        this.n = n;  
        System.out.println("Put: " + n);  
    }  
}
```

```
class Producer implements Runnable {
    Q q;
    Producer(Q q) {
        this.q = q;
        new Thread(this, "Producer").start();
    }
    public void run() {
        int i = 0;
        while (true) {
            q.put(i++);
        }
    }
}
```

```
class Consumer implements Runnable {
    Q q;
    Consumer(Q q) {
        this.q = q;
        new Thread(this, "Consumer").start();
    }
    public void run() {
        while (true) {
            q.get();
        }
    }
}
```

```
class PC {  
    public static void main(String args[]) {  
        Q q = new Q();  
        new Producer(q);  
        new Consumer(q);  
        System.out.println("Press Control-C to stop.");  
    }  
}
```

PRODUCER-CONSUMER PROBLEM

- ◉ Correct Solution


```
class Q {  
    int n;  
    boolean valueSet = false;  
    synchronized int get() {  
        while (!valueSet)  
            try {  
                wait();  
            } catch (InterruptedException e) {  
                e.printStackTrace();  
            }  
        System.out.println("Got: " + n);  
        valueSet = false;  
        notify();  
        return n;  
    }  
}
```

```
synchronized void put(int n) {  
    while (valueSet)  
        try {  
            wait();  
        } catch (InterruptedException e) {  
            e.printStackTrace();  
        }  
    this.n = n;  
    valueSet = true;  
    System.out.println("Put: " + n);  
    notify();  
}  
}
```

```
class Producer implements Runnable {
    Q q;
    Producer(Q q) {
        this.q = q;
        new Thread(this, "Producer").start();
    }
    public void run() {
        int i = 0;
        while (true) {
            q.put(i++);
        }
    }
}
```

```
class Consumer implements Runnable {
    Q q;
    Consumer(Q q) {
        this.q = q;
        new Thread(this, "Consumer").start();
    }
    public void run() {
        while (true) {
            q.get();
        }
    }
}
```

```
class PCCorrectDemo {  
    public static void main(String args[]) {  
        Q q = new Q();  
        new Producer(q);  
        new Consumer(q);  
        System.out.println("Press Control-C to stop.");  
    }  
}
```

THANK YOU